

# Introduction to change point detection

---

Euan T. McGonigle

12th September

University of Southampton

# Session Plan

First part:

- Learn about what change point detection is.
- Learn what the common change point methods are.
- Learn about some software we can use.

Second part:

- Write code and use software to try out change point algorithms.
- Compare methods and see the pros/cons.
- Perform change point analysis on some environmental data sets.

All materials for today are available on my Github page:  
[github.com/euanmcgonigle/IntroToCPD](https://github.com/euanmcgonigle/IntroToCPD).

# Software we'll use

I'll be using R for the session, but if you are more comfortable then please use Python: both have everything we need for what we'll cover today.

If you want to use R, you'll need:


- changepoint
- mosum
- EnvCpt

If you want to use Python, you'll need:

- ruptures
- mosum

# Introduction

- Time series are now recorded in higher volumes over longer periods.
- Often, dynamics of time series are assumed stationary – unrealistic!
- Instead, we can develop nonstationary time series models.
- Could be with respect to mean level, spectral density, correlation, ...
- Examples: climatology, finance, neuroscience, ...
- 1 approach: piecewise stationarity via change point detection.



# Introduction

- We might want to take a change point approach to model things like:
  1. air quality,
  2. frequencies of earthquakes,
  3. global sea temperatures,
  4. sentiment of text data like tweets.
- We'll focus on some common time series models to illustrate key ideas, but these techniques can be readily adapted to other more complicated models.

# Change point detection (CPD)


Change points: abrupt changes in the statistical properties of the time series. Also known as:

- data segmentation,
- break points,
- structural breaks,
- data shift/drift, ...

We'll look at some of the most well-studied settings, and see how some simple ideas can be used to build powerful CPD algorithms.


# Change point detection (CPD)

Change points could occur in any properties of the data we want to model, such as the mean, dependence, and trend:



# More complicated models


How many change points below? And where?





# More complicated models


How many change points below? And where?



A time series is split into stationary segments defined by the change points. We want to:


1. **Test:** Have changes occurred? If so, how many?
2. **Localise:** Where have the changes happened?
3. **Theoretically justify:** Statistical guarantee on number of changes and accuracy of estimated locations.
4. **Analyse post-segmentation:** Estimate parameters, quantify uncertainty, forecasting, causality?

## Example application: global sea temperature anomalies



## Example application: earthquake occurrences

Monthly number of earthquakes in the Groningen oil field:



# Why CPD?

- The simplest possible nonstationary model.
- Very interpretable.
- Tends to lead to computationally tractable and exact solutions.
- Often we can tie change points to physical phenomena e.g. stock market crash, policy interventions, human-made emissions.
- Data-driven, “agnostic” intervention analysis.

## Simple model – change in mean

Let's look at the canonical change in mean problem:

$$X_t = f_t + \varepsilon_t = \sum_{j=0}^q \mu_j \cdot \mathbb{I}\{\theta_j + 1 \leq t \leq \theta_{j+1}\} + \varepsilon_t, \quad t = 1, \dots, n.$$

- $f_t$  is the signal – piecewise constant mean function.
- $f_t$  has jumps at the  $q$  change points  $\theta_1, \dots, \theta_q$ .
- $\varepsilon_t$  is a zero-mean noise term.

Note: for a general setting:

$$X_t = \sum_{j=0}^q X_t^{(j)} \cdot \mathbb{I}\{\theta_j + 1 \leq t \leq \theta_{j+1}\},$$

where distributions of  $X^{(j)}$  and  $X^{(j+1)}$  differ somehow.

## Aside – online vs offline

Two main settings for CPD are **offline** and **online**.



Online:

- Processes the data as it arrives, or in batches.
- Main aim is to quickly detect new changes.
- Applications include process control monitoring, fault detection.

Offline:

- Processes all the data in one go, after it has all been collected.
- Main aim is accurate detection of changes.
- Applications include genome analysis.

## Aside – online vs offline





# Simple model – change in mean

What do we want to do?

- Estimate the number of change points  $q$ .
- Estimate the change point locations  $\theta_1, \dots, \theta_q$ .

How do we do that?

Many approaches, 2 of which are:

- Localised testing using a test statistic.
- Cost function based on data fit.

# Localised testing

- Scan data for candidate change point estimators.
- Typically involve computing a test statistic of the form  $\hat{\sigma}_{s,e}^{-1}|\mathcal{T}_{s,k,e}|$ , where

$$\mathcal{T}_{s,k,e} = \sqrt{\frac{(k-s)(e-k)}{e-s}} \left( \frac{1}{k-s} \sum_{t=s+1}^k X_t - \frac{1}{e-k} \sum_{t=k+1}^e X_t \right),$$

and  $\hat{\sigma}_{s,e}$  is a estimator of  $\sigma_{s,e}$ , a measure of variability of  $\{X_t\}_{t=s}^e$ .

- Compute  $\mathcal{T}_{s,k,e}$  for all  $k$  over a range of intervals to find the most likely change point locations.
- Often compare  $\hat{\sigma}_{s,e}^{-1}|\mathcal{T}_{s,k,e}|$  to a threshold  $D$  to test for changes.

## CUSUM approaches: idea

- Simplest example is the cumulative sum (CUSUM):

$$\mathcal{T}_{0,k,n} = \sqrt{\frac{k(n-k)}{n}} \left( \frac{1}{k} \sum_{t=1}^k X_t - \frac{1}{n-k} \sum_{t=k+1}^n X_t \right).$$


- To find a single change  $\theta_1$  we compute  $\mathcal{T}_{0,k,n}$  for all  $1 \leq k \leq n-1$ , and set

$$\hat{\theta}_1 = \arg \max_{1 \leq k \leq n-1} |\mathcal{T}_{0,k,n}|$$


to be the change point estimator.

- We declare  $\hat{\theta}_1$  as a change if  $\hat{\sigma}^{-1} |\mathcal{T}_{0,\hat{\theta}_1,n}| > D$ .
- Many CPD methods start with the CUSUM as the basic idea, but use fancier ideas.

# CUSUM approaches: single change




# CUSUM approaches: multiple changes



## CUSUM approaches: extensions

- It's easy to extend the global CUSUM to look for multiple change points.
- In Binary Segmentation (BS), after calculating  $\hat{\theta}_1$ , we then perform 2 CUSUMs on the intervals  $(1, \hat{\theta}_1)$  and  $(\hat{\theta}_1 + 1, n)$ , and rinse and repeat.
- Wild binary segmentation extends this by using BS on a set of randomly generated intervals.
- Many other state-of-the-art methods devise sophisticated “interval generation” techniques, upon which CUSUMs are performed.

# CUSUM approaches: smart interval generation



There are many other CPD methods that use localised testing.

- The **MOSUM** approach uses moving windows of size  $G$  to the left and right of candidate change  $k$ .
- This allows it to simultaneously find multiple changes.
- For window size  $G$ , the MOSUM detector  $T_G(k)$  is given by


$$T_G(k) = T_{k-G,k,k+G} = \frac{1}{\sqrt{2G}} \left( \sum_{t=k+1}^{k+G} X_t - \sum_{t=k-G+1}^k X_t \right)$$

for  $k = G, \dots, n - G$ .

- Changes declared as all local maximisers with  $\hat{\sigma}^{-1} |T_G(k)| > D$ .



# MOSUM approaches: multiple change points



# MOSUM approaches: extensions

- Again, this is a simple idea that can be extended into more sophisticated algorithms.
- Uncertainty quantification on change locations via bootstrap.
- Instead of using  $G$  data points on each side of the window, can use asymmetric sizes.
- A multiscale version of MOSUM using multiple bandwidths helps to detect changes of different sizes and different segment lengths.
- Can easily be generalised to multivariate data and other types of change.

## MOSUM approaches: extensions

- Again, this is a simple idea that can be extended into more sophisticated algorithms.
- Uncertainty quantification on change locations via bootstrap.
- Instead of using  $G$  data points on each side of the window, can use asymmetric sizes.
- A multiscale version of MOSUM using multiple bandwidths helps to detect changes of different sizes and different segment lengths.
- Can easily be generalised to multivariate data and other types of change.

# Cost function approaches

Another common approach is using cost functions.

- A cost function  $\mathcal{C}(X_{s:e})$  measures how well (badly) the model fits the data on the segment  $\{X_t\}_{t=s}^e$ .
- For 1 change point, we can measure the cost of fitting a change point at time  $k$  as  $\mathcal{C}(X_{1:k}) + \mathcal{C}(X_{(k+1):n})$ .
- For finding  $q$  change points, we can minimise the equation

$$\sum_{j=1}^{q+1} \mathcal{C}(X_{(\theta_{j-1}+1):\theta_j}) + f(q),$$

where  $f(q)$  is a penalty to avoid overfitting.

- Twice the negative log likelihood is a commonly used cost function.



## Cost function approaches: PELT

- Trying to calculate this minimum exhaustively is not feasible.
- Thankfully it can be calculated recursively: the minimal cost for  $X_{1:s}$  can be expressed in terms of that of  $X_{1:t}$  for  $t < s$ .
- This is still complexity  $\mathcal{O}(n^2)$ : the Pruned Exact Linear Time (PELT) method improves this to expected complexity  $\mathcal{O}(n)$ .
- Achieves this by pruning the set of possible change point locations in a clever way, using properties of the cost function.
- For change in mean, equivalent to least-squares estimation.

# Nonparametric approaches

What if we're not sure what type of parameters/model we want to pick?

We can miss changes if we don't know what we're looking for, or falsely detect changes, if we assume the wrong model:



In this case, we can use nonparametric methods.

# Nonparametric approaches

- We can do a kernel version of the least-squares/cost function approach: for given  $q$ , minimise


$$L(q) = \frac{1}{n} \sum_{i=1}^n k(X_i, X_i) - \frac{1}{n} \sum_{j=1}^q \left[ \frac{1}{\theta_j - \theta_{j-1}} \sum_{k=\theta_{j-1}+1}^{\theta_j} \sum_{\ell=\theta_{j-1}+1}^{\theta_j} k(X_k, X_\ell) \right],$$

where  $k(\cdot, \cdot)$  is a kernel function.

- Many choices of kernel: most common is the Gaussian  $k(x, y) = \exp(-||x - y||^2 / (2h^2))$ .
- If  $k(x, y) = xy$ , this becomes the same as the PELT method for changes in mean.
- We are basically just turning the problem into mean change point detection via kernels.

# Nonparametric approaches


Example approach: a single test statistic/cost function can detect multiple types of changes.





# Model validation/checking assumptions

- Often we want to validate the assumptions of our model, e.g. normality, independence of data.
- Everything becomes more difficult in the presence of change points:



- Can use basic diagnostics like checking residuals, e.g., QQ or ACF plots, with caution.
- We can use model selection approaches to compare competing change point models, e.g., with information criterion.

# Multivariate data

All of the above discussion considered univariate data. What about multivariate data?

- Each variable of the data might be unconnected to the rest: univariate methods could work well.
- There might be structure/dependence across variables. For example,
  1. Changes occur at the same time across all variables,
  2. changes occur only in a subset of variables,
  3. changes in some variables cause changes in others at a later time.
- Many possible other scenarios: different types of change, different noise levels, ...
- Generally, we can take ideas from univariate methods and extend them to multivariate data.

## Aside: multivariate change point software

Many packages doing multivariate change point detection.

In R:

- `ecp`: nonparametric CPD using “energy” distance.
- `InspectChangepoint`: high-dimensional mean CPD using projections.
- `fastcpd`: many settings, including several from `ruptures`.

In Python:

- `ruptures`: already has multivariate functionality.
- `gchangeoint`: graph-based nonparametric CPD.
- `changeforest`: nonparametric CPD via random forests.

## Complications/considerations

- Speed of method: sequential nature means efficient implementations often possible.
- High-dimensionality: how do we aggregate information/find relevant structure?
- Outliers: “instantaneous” change points? Many methods not robust.
- Threshold selection: how to balance over/under detection?
- Missing data: how to impute when there are change points?
- Autocorrelation: harder to distinguish changes from noise.

## Aside: some alternative approaches

Whilst I'm contractually obliged to advertise CPD, it's only 1 of many approaches:

- Mathematical/physical model: something that already captures the evolution of behaviour of interest.
- Local stationarity: statistical properties evolve slowly.
- Just ignore it (!): for some applications, e.g. forecasting, it might be better to just fit a stationary model.
- Data pre-processing: remove nonstationarity first, e.g. via a transformation.

## Some things to think about

- What types of time series might contain change points?
- What types of change points might you expect to see in the type of data you look at.
- What parameters you might model that have nonstationary behaviour.
- Benefits and drawbacks to a change point approach.
- Homogeneity vs inhomogeneity – over space, time, ...

## Further reading

- Eichinger, B. and Kirch, C. (2018). A mosum procedure for the estimation of multiple random change points. *Bernoulli*, 24(1):526–564.
- Fearnhead, P. and Fryzlewicz, P. (2022). Detecting a single change-point. *arXiv preprint arXiv:2210.07066*.
- Fryzlewicz, P. (2014). Wild binary segmentation for multiple change-point detection. *The Annals of Statistics*, 42(6):2243–2281.
- Killick, R. and Eckley, I. A. (2014). changepoint: An r package for changepoint analysis. *Journal of Statistical Software*, 58:1–19.
- Killick, R., Fearnhead, P., and Eckley, I. A. (2012). Optimal detection of changepoints with a linear computational cost. *Journal of the American Statistical Association*, 107(500):1590–1598.
- Meier, A., Kirch, C., and Cho, H. (2021). mosum: A package for moving sums in change point analysis. *Journal of Statistical Software*, 97(8):1–42.
- Truong, C., Oudre, L., and Vayatis, N. (2018). ruptures: change point detection in python. *arXiv preprint arXiv:1801.00826*.